

Aarav Sharma

San Francisco, CA | aarav.sharma@example.com | +1 (415) 555-0182 | linkedin.com/in/aaravsharma-ml | github.com/aaravsharma

PROFESSIONAL SUMMARY

Accomplished Lead Machine Learning Engineer with 5 years of experience building scalable deep learning systems, multi-modal transformer models, and high-throughput REST inference APIs. Proven track record deploying containerized models in production on AWS with sub-50ms latency.

PROFESSIONAL EXPERIENCE

DeepVision AI — Senior Machine Learning Engineer (2022 - Present)

- Architected and trained multi-modal transformer models in PyTorch, improving contextual search accuracy by 34% and reducing inference latency by 42%.
- Built high-throughput asynchronous inference microservices using FastAPI, Redis caching, and Docker deployed on AWS ECS.
- Created robust feature engineering pipelines querying multi-terabyte PostgreSQL and Snowflake data warehouses.
- Led a team of 4 junior ML engineers and established MLOps best practices with MLflow and GitHub Actions.

Apex Data Labs — Machine Learning Engineer (2020 - 2022)

- Developed predictive classification and time-series forecasting models using Scikit-learn and TensorFlow.
- Implemented automated CI/CD model evaluation and data validation pipelines using Git, Docker, and MLflow.
- Optimized SQL ETL queries reducing data pipeline runtimes from 4 hours to 45 minutes.

EDUCATION

Stanford University — Master of Science in Computer Science (AI Track), 2018 - 2020 (GPA: 3.9/4.0)

IIT Bombay — Bachelor of Technology in Computer Science, 2014 - 2018 (GPA: 3.85/4.0)

TECHNICAL SKILLS

- Programming Languages: Python, SQL, C++, Bash
- Frameworks: PyTorch, TensorFlow, FastAPI, Scikit-learn, HuggingFace Transformers, LangChain
- Databases & Data: PostgreSQL, Redis, Snowflake, BigQuery
- AI/ML: Deep Learning, NLP, RAG, Computer Vision, Transformers, LLM Fine-tuning, MLOps
- Tools & Cloud: Docker, Kubernetes, AWS (EC2, S3, ECS, Lambda), Git, Linux, CI/CD
- Soft Skills: Technical Leadership, Cross-functional Collaboration, Mentorship, Problem Solving

KEY PROJECTS

- Enterprise RAG Platform: Built an enterprise retrieval-augmented generation engine using LangChain, PyTorch, and Milvus vector search serving 10,000+ daily queries with 98% precision.
- High-Speed Vision Classifier: Deployed quantized ResNet and ViT models on AWS Lambda with sub-50ms p95 latency.

CERTIFICATIONS

- AWS Certified Machine Learning - Specialty (2023)